

MNIST Dataset Overview

This example is using MNIST handwritten digits. The dataset contains 60,000 examples for training and 10,000 examples for testing. The digits have been size-normalized and centered in a fixed-size image (28x28 pixels) with values from 0 to 255.

```
In [1]:
from __future__ import absolute_import, division, print_function

import tensorflow as tf
from tensorflow.keras import Model, layers
import numpy as np

In [2]:
# MNIST dataset parameters.
num_classes = 10 # total classes (0-9 digits).

# Training parameters.
learning_rate = 0.001
training_steps = 25
batch_size = 64
display_step = 10

# Network parameters.
conv1_filters = 32 # number of filters for 1st conv layer.
conv2_filters = 64 # number of filters for 2nd conv layer.
fc1_units = 1024 # number of neurons for 1st fully-connected layer.

In [3]:
# Prepare MNIST data.
from tensorflow.keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
# Convert to float32.
x_train, x_test = np.array(x_train, np.float32), np.array(x_test, np.float32)
# Normalize images value from [0, 255] to [0, 1].
x_train, x_test = x_train / 255., x_test / 255.

In [4]:
# Use tf.data API to shuffle and batch data.
train_data = tf.data.Dataset.from_tensor_slices((x_train, y_train))
train_data = train_data.repeat().shuffle(5000).batch(batch_size).prefetch(1
)
```

```

In [5]:
# Create TF Model.
class ConvNet(Model):
    # Set layers.
    def __init__(self):
        super(ConvNet, self).__init__()
        # Convolution Layer with 32 filters and a kernel size of 5.
        self.conv1 = layers.Conv2D(32, kernel_size=5, activation=tf.nn.relu
)
        # Max Pooling (down-sampling) with kernel size of 2 and strides of
2.
        self.maxpool1 = layers.MaxPool2D(2, strides=2)

        # Convolution Layer with 64 filters and a kernel size of 3.
        self.conv2 = layers.Conv2D(64, kernel_size=3, activation=tf.nn.relu
)
        # Max Pooling (down-sampling) with kernel size of 2 and strides of
2.
        self.maxpool2 = layers.MaxPool2D(2, strides=2)

        # Flatten the data to a 1-D vector for the fully connected layer.
        self.flatten = layers.Flatten()

        # Fully connected layer.
        self.fc1 = layers.Dense(1024)
        # Apply Dropout (if is_training is False, dropout is not applied).
        self.dropout = layers.Dropout(rate=0.5)

        # Output layer, class prediction.
        self.out = layers.Dense(num_classes)

    # Set forward pass.
    def call(self, x, is_training=False):
        x = tf.reshape(x, [-1, 28, 28, 1])
        x = self.conv1(x)
        x = self.maxpool1(x)
        x = self.conv2(x)
        x = self.maxpool2(x)
        x = self.flatten(x)
        x = self.fc1(x)
        x = self.dropout(x, training=is_training)
        x = self.out(x)
        if not is_training:
            # tf cross entropy expect logits without softmax, so only
            # apply softmax when not training.
            x = tf.nn.softmax(x)
        return x

```

```
# Build neural network model.
```

```
conv_net = ConvNet()
```

```
In [6]:
```

```
# Cross-Entropy Loss.
```

```
# Note that this will apply 'softmax' to the logits.
```

```
def cross_entropy_loss(x, y):
```

```
    # Convert labels to int 64 for tf cross-entropy function.
```

```
    y = tf.cast(y, tf.int64)
```

```
    # Apply softmax to logits and compute cross-entropy.
```

```
    loss = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y, logits=
x)
```

```
    # Average loss across the batch.
```

```
    return tf.reduce_mean(loss)
```

```
# Accuracy metric.
```

```
def accuracy(y_pred, y_true):
```

```
    # Predicted class is the index of highest score in prediction vector (i
.e. argmax).
```

```
    correct_prediction = tf.equal(tf.argmax(y_pred, 1), tf.cast(y_true, tf.
int64))
```

```
    return tf.reduce_mean(tf.cast(correct_prediction, tf.float32), axis=-1)
```

```
# Stochastic gradient descent optimizer.
```

```
optimizer = tf.optimizers.Adam(learning_rate)
```

```
In [7]:
```

```
# Optimization process.
```

```
def run_optimization(x, y):
```

```
    # Wrap computation inside a GradientTape for automatic differentiation.
```

```
    with tf.GradientTape() as g:
```

```
        # Forward pass.
```

```
        pred = conv_net(x, is_training=True)
```

```
        # Compute loss.
```

```
        loss = cross_entropy_loss(pred, y)
```

```
    # Variables to update, i.e. trainable variables.
```

```
    trainable_variables = conv_net.trainable_variables
```

```
    # Compute gradients.
```

```
    gradients = g.gradient(loss, trainable_variables)
```

```
# Update W and b following gradients.
optimizer.apply_gradients(zip(gradients, trainable_variables))
```

In [8]:

```
# Run training for the given number of steps.
for step, (batch_x, batch_y) in enumerate(train_data.take(training_steps),
1):
    # Run the optimization to update W and b values.
    run_optimization(batch_x, batch_y)

    if step % display_step == 0:
        pred = conv_net(batch_x)
        loss = cross_entropy_loss(pred, batch_y)
        acc = accuracy(pred, batch_y)
        print("step: %i, loss: %f, accuracy: %f" % (step, loss, acc))
```

In [9]:

```
# Test model on validation set.
pred = conv_net(x_test)
print("Test Accuracy: %f" % accuracy(pred, y_test))
```

In [10]:

```
# Visualize predictions.
import matplotlib.pyplot as plt
```

In [11]:

```
# Predict 5 images from validation set.
n_images = 5
test_images = x_test[:n_images]
predictions = conv_net(test_images)

# Display image and model prediction.
for i in range(n_images):
    plt.imshow(np.reshape(test_images[i], [28, 28]), cmap='gray')
    plt.show()
    print("Model prediction: %i" % np.argmax(predictions.numpy()[i]))
```